

Master of Science HES-SO in Engineering

Orientation : Computer science (CS)

Cloud-Native IoT Reference Architecture with Arm SystemReady

Made by
Xavier Clivaz

Under the direction of
Prof. Fabien Vannel
Hepia/inIT/CoRES

Information about this report

Contact Information

Author: Xavier Clivaz
Master Student
HES-SO//Master
Switzerland
Email: xavierclivaz@hotmail.com

Declaration of honor

I, undersigned, Xavier Clivaz, hereby declare that the work submitted is the result of a personal work. I certify that I have not resorted to plagiarism or other forms of fraud. All sources of information used and the author quotes were clearly mentioned.

Place, date: Sion, January 29, 2024

Signature:



Acknowledgements

This master's thesis would not have been written over five months without the support of a number of people. It's time to thank all these contributors, whether they are closely involved in the project or are far removed from the digital world.

- First of all, I would like to thank my family and especially my parents, without whom I wouldn't be able to fulfil all my dreams. So I'm always able to fulfil my potential in my daily life. It's the strength of spirit I've inherited that will always enable me to move forward despite the obstacles I encounter.
- Secondly, I would like to thank my supervisor, Mr Fabien Vannel, who was always at my complete disposal. He was able to provide me with invaluable advice throughout the work to achieve the desired result.
- Finally, I would like to thank everyone at [56K.Cloud](#). They are a great team, bringing joy and good humour to the organisation. What's more, they were able to contribute their knowledge when obstacles were encountered.

Abstract

Key words: Arm SystemReady, AWS, Cloud-Native, IoT, reference architecture

Contents

Acknowledgements	v
Abstract	vii
Contents	viii
List of Figures	x
List of Listings	x
1 Introduction	1
1.1 56K.Cloud	1
1.2 Problem	2
1.3 Objectives	3
1.4 Project plan	3
1.5 Research methodology	4
1.6 Structure of this report	5
2 Project methodology	7
2.1 Agile methodology	8
2.2 KanBan methodology	10
2.3 Project management tools	11
2.4 Formations	13
3 Analysis	15
3.1 Definitions	16
3.2 Problem formulation	18
3.3 Literature search	19
3.4 History of the Cloud-Native approach	19
3.5 Integrating cloud computing with IoT embedded systems	21
3.6 IoT cloud platforms	25
3.7 Cloud infrastructure tools	27
3.8 Arm SystemReady	31
4 Design	33
4.1 Section 1	34
5 Implementation	35
5.1 Section 1	36

6	Validation	37
6.1	Section 1	38
7	Conclusions	39
7.1	Project summary	39
7.2	Comparison with the initial objectives	39
7.3	Encountered difficulties	39
7.4	Future perspectives	39
A	Plannings	41
A.1	Forward planning	41
A.2	Effective planning	43
	Bibliography	45
	Glossary	51
	Acronyms	53

List of Figures

1.1	56K.Cloud company logo [1]	1
1.2	The problem of linking the hardware to the cloud infrastructure [2]	2
1.3	Overview of objectives	3
1.4	Project plan	4
2.1	Sequential and iterative process [4]	8
2.2	Scrum roles	9
2.3	Scrum life cycle	10
2.4	Example of a virtual KanBan table [10]	11
2.5	Example of version management with two branches	11
2.6	D evelopment and O perations (DevOps) method	12
3.1	Comparison of different service models [12]	17
3.2	Google trends (01.01.2006 until 22.05.2016) of term Cloud-Native [17]	19
3.3	Traditional server (left) and virtualisation (right)	20
3.4	Model for integrating cloud computing into an embedded system [27]	22
3.5	Cloud Customer Reference Architecture for IoT [29]	23
3.6	Comparison of IoT cloud platform providers (sources from providers)	27
3.7	Overview of AWS CDK [70] deployment [70]	28
3.8	Overview of Terraform deployment [72]	29
3.9	Overview of Pulumi deployment [73]	30

List of Listings

1 | Introduction

For a long time, embedded systems remained isolated from each other. Now it's time to bring them together with the world of the [cloud](#). Hardware devices have always been limited in terms of resources. Designs have allowed them to evolve according to their use case. However, physical limits have been observed. This is how the [cloud computing](#) came into being, partly in response to these difficulties.

The [cloud](#) provides a huge number of advantages. It frees up hardware resources to optimise the operation of embedded systems. It is likely that these systems will not be able to perform complex calculations. The [cloud computing](#) has the capacity to do this. The storage problem is another constraint. The [cloud computing](#) offers almost unlimited storage. It doesn't stop there. It is capable of offering constant scalability to infrastructures. It is made up of different independent services that can be easily connected together. There is vertical scaling, which means increasing the size of a resource, and horizontal scaling, which means increasing the number of resources.

More and more engineers have decided to take an interest in this area in order to interconnect their embedded systems with the [cloud](#). This practice is an integral part of the [Internet of Things](#). Due to increasingly large projects, not to mention artificial intelligence, they want to use [cloud computing](#) environments in order to no longer be constrained by resources. To make it easier to use a [cloud infrastructure](#), companies have decided to offer [cloud](#) platforms comprising a multitude of services for all types of use ([AWS](#), Microsoft Azure, etc). It's also a way of letting engineers work on their business rather than worrying about maintaining and securing these infrastructures.

1.1 56K.Cloud

[56K.Cloud](#) is a company based in Sion (Valais) since 2018. It offers various services through the [cloud](#). It provides solutions for businesses that want to improve their processes while reducing their costs. It is also a consultancy firm that seeks to promote the [cloud](#) to customers who find it difficult to understand certain concepts in this rapidly expanding digital world. Collaborations with a number of partners are also an asset, enabling us to pool a wide range of skills to provide customers with a complete product. In order to share its vision of the [cloud](#) more widely, it has additional premises in Winterthur (Zurich).



Figure 1.1 [56K.Cloud](#) company logo [1]

The idea of a [cloud](#) solution is to move an on-premise infrastructure to the [cloud](#). It's a way of reducing costs by only paying for what you use (pay-as-you-go), and by having something that can be easily scalable. It's also a good way of increasing the speed of processes by using optimal resources for each application.

In terms of consultancy, 56K.Cloud introduces Cloud-Native technology to customers who want to find out more. It also covers the subject of DevOps and containerisation.

1.2 Problem

While 56K.Cloud mainly offers cloud solutions, this company also works in the Internet of Things (IoT) sector. Recently, it quickly noticed a major problem between embedded systems engineers and those from the world of cloud. More and more engineers are asking for devices to be able to link easily and quickly to a cloud infrastructure. However, there is currently no reference architecture that enables an infrastructure to be deployed and an embedded system to be provisioned directly.

This is a real problem because of the time needed and the difficulty for engineers to link hardware to a cloud infrastructure. Because of the two distinct professions, a person coming from the cloud will need time to understand and set up an operating system capable of connecting to an infrastructure. And vice-versa, a person coming from the hardware world will need to familiarise themselves with a cloud infrastructure, not to mention setting it up while using the best practices of the Cloud-Native.




Figure 1.2 The problem of linking the hardware to the cloud infrastructure [2]

The need to find a solution arose when engineers started asking electronic board manufacturers to design equipment offering services for linking to a cloud infrastructure. Since these manufacturers are not in this discipline, they do not want to spend time developing these types of solutions. It is therefore possible to identify the main players affected. There are the hardware engineers who are unable to meet the new needs of customers. Cloud engineers, who want to do the job for them, have to devote part of their time to introducing embedded concepts. Learning something new can quickly become time-consuming.

Nevertheless, a few solutions are gradually emerging. However, these remain very closed. There is no open source reference architecture capable of satisfying a wide range of products. These solutions therefore remain highly proprietary.

1.3 Objectives

The primary aim of this work is to develop a reference architecture enabling the deployment of a [cloud infrastructure](#) with the [provisioning](#) of a fleet of embedded systems. The idea is to be able to automatically provision devices to the infrastructure when they are first started up. The infrastructure must contain the essential minimum of components to make the architecture as universal as possible. Since [56K.Cloud](#) is a partner of [Arm](#), a processor manufacturer, the embedded systems must be equipped with this. It should be added that [Arm](#) puts certifications on their components such as SystemReady. SystemReady program guarantees that an operating system and the following software layers can function correctly in their processors. The embedded systems used in this project must be certified SystemReady. In addition, the infrastructure must be deployed on the [AWS](#) platform. This is the [cloud](#) provider that [56K.Cloud](#) works with.

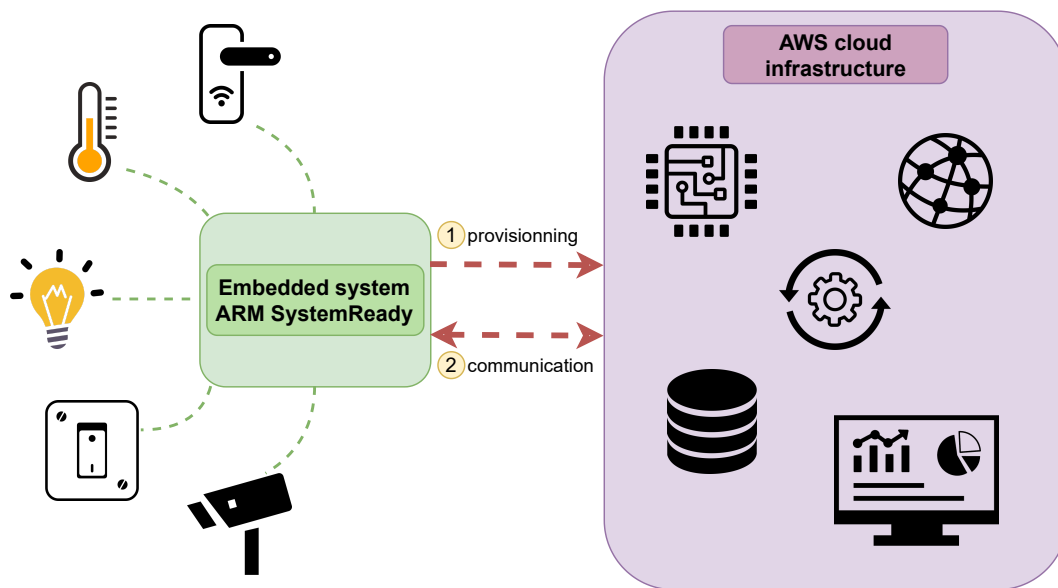


Figure 1.3 Overview of objectives

The aim of this project is to be open source so that a community can form around it. It should enable embedded systems and [cloud](#) engineers to use this reference architecture easily. They need to be able to focus on their final product. The best practices of [Cloud-Native](#) must be used for the development of this project. In order to guarantee a sound architecture, the use of [Continuous Intergration \(CI\)](#) and [Continuous Delivery \(CD\)](#) tools are essential. In addition, to ensure compatibility, the architecture must be functional on different embedded systems with a [Arm](#) processor and certified SystemReady.

Secondly, once the reference architecture is functional, a demonstration project must be based on it. This is a proof of concept. The general idea is to deploy a [cloud infrastructure](#) on [AWS](#) and have various [cloud](#) services automatically set up to interact with the embedded system. Data will then transit between these two worlds and it must be viewable from an interface.

1.4 Project plan

Figure 1.4 shows the project plan in the form of [work packages](#).

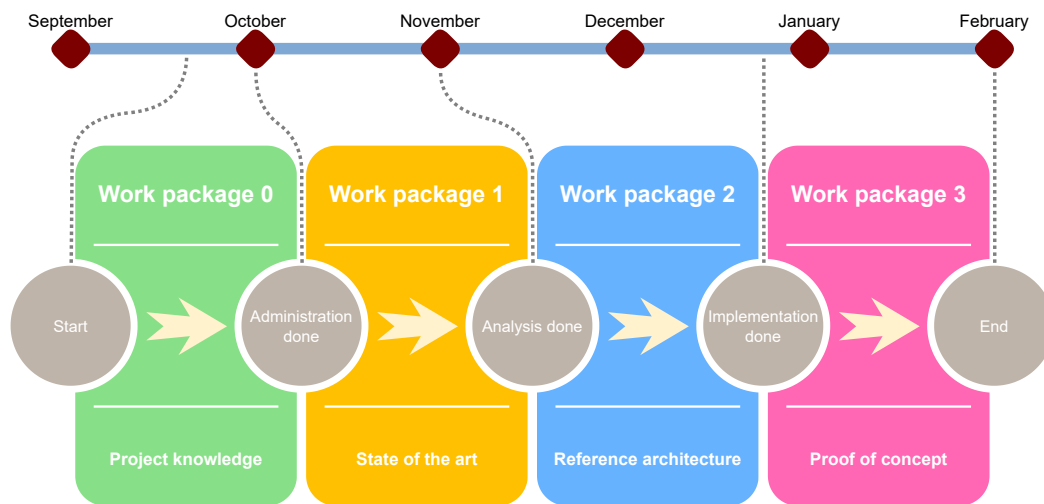


Figure 1.4 Project plan

The first [work package](#) contains all the administration. It involves learning about the project, being aware of the issues and clearly defining the specifications. This takes around two weeks.

About a month is spent researching existing solutions that are closest to the project. This involves drawing up a state of the art by researching scientific articles and other documentation.

The third [work package](#) concerns the reference architecture. Almost two months are devoted to this. This implementation part will be the final product, which will be delivered as open source on a shared repository.

In order to validate this architecture, a proof of concept will be carried out. This is the final [work package](#). It will take just over a month. Tests will be carried out to prove that it works properly.

1.5 Research methodology

The research methodology used in this thesis was first to draw up a state of the art on existing solutions that are as close as possible to this project. An overview of the reference architecture was then drawn up with a view to implementing it more easily. A choice of development tools was also made. At the same time, the implementation was designed. Finally, the results were validated by means of a proof of concept.

The agile methodology, known as [Scrum](#), was used throughout the project. It allows for iterative project management. Weekly meetings were arranged with the professor in charge of the project and, optionally, with some of the company's staff. To enable the project to progress efficiently, the [KanBan](#) method was used. This consists of virtual cards, each referring to a task. These cards are categorised in a table to determine the progress of each task. It's also a way of dividing the project into smaller parts so that time can be better estimated over the course of the project.

1.6 Structure of this report

The 2 chapter ([Project methodology](#)) contains a number of project methodologies for effective project management. The tools used are also described.

The 3 chapter ([Analysis](#)) contains several definitions and the state of the art. These are definitions of important terms that make up the project. The state of the art concerns the existing solutions that are closest to this project.

The 4 chapter ([Design](#)) contains an overview of the implementation as well as ...
To be completed

The 5 chapter ([Implementation](#)) contains ... **To be completed**

The 6 chapter ([Validation](#)) contains the analysis and validation of the results of the implementation.

The 7 chapter ([Conclusions](#)) concludes the thesis by outlining the current state of the research, the problems encountered and future steps.

2 | Project methodology

Project methodology consists of describing the means used to manage a project. In this chapter, the concept of agility is described in order to explain how the iterative process works. In addition, the [KanBan](#) methodology is described in detail to explain the time aspect of the tasks to be carried out. Finally, the various tools used are discussed.

Contents

2.1	Agile methodology	8
2.1.1	Scrum roles	9
2.1.2	Scrum process	9
2.2	KanBan methodology	10
2.3	Project management tools	11
2.3.1	Version management	11
2.3.2	Continuous Intergration (CI) and Continuous Delivery (CD)	12
2.3.3	Work packages	13
2.4	Formations	13

2.1 Agile methodology

Agile methodology is the main project management method used in this work. It enables the project to be managed from the specifications to the final product. The idea for this method was conceived in the 70s or even before [3]. The idea was to change the development process in software engineering by using iterative techniques. Traditional methods worked with sequential techniques, often called Waterfall. Figure 2.1 shows the difference between these two approaches.

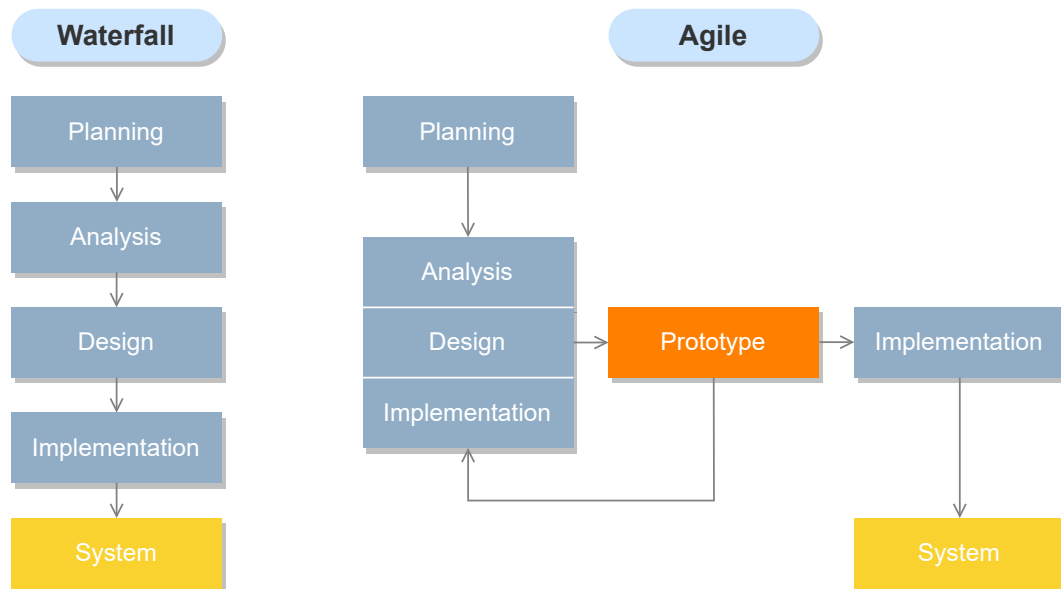


Figure 2.1 Sequential and iterative process [4]

The Waterfall methodology is considered cumbersome. There are deliverables after each phase. Approval is required before moving on to the next phase. It is difficult to go backwards (by moving up the phases). Professors explained this very well during a Masters course given at the HES-SO [4]. They added that this method is best used in very large projects where it is difficult to split teams into small groups to work iteratively.

These same teachers [4] described the agile methodology as follows:

"AGILE" is about values and principles, not practices, but many practices support them. It's not about doing agile, it's about being agile.

Several agile process frameworks were created before the 2000s. Examples include Scrum, XP, RUP, etc. Between the 11th and the 13th of February 2001, 17 people from these different frameworks met to find an alternative to cumbersome, documentation-driven software development processes. They created the agile software development manifesto [5] :

*Individuals and interactions over processes and tools
Working software over comprehensive documentation
Customer collaboration over contract negotiation
Responding to change over following a plan*

The central principle of agile is to have very rapid iterations to build the first business values. The customer is at the centre of the process thanks to frequent communication with the development team [4]. The process framework chosen for this project is **Scrum**. It is very well suited to small projects with small teams. It enables the first software prototypes to be delivered quickly. It adapts easily to changes, unlike a sequential method. It is also based on experience. The idea is to learn continuously through iterations. This is an important point in terms of learning from new developments in the project. This is one of the most popular agile methods. This approach was launched in 1995 by Ken Schwaber and Jeff Sutherland [6].

2.1.1 Scrum roles

Figure 2.2 shows the different roles in a **Scrum** team.

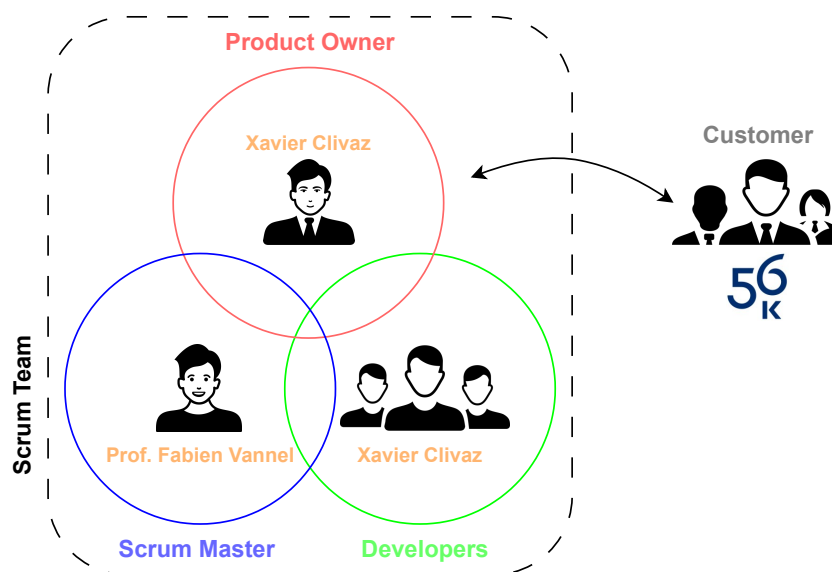


Figure 2.2 Scrum roles

A **Scrum** guide has been written by Jeff Sutherland and Ken Schwaber to explain the rules [7]. First of all, there is the **Product Owner (PO)**. This is the customer's spokesperson. It is he who defines the product's functionalities. They are recorded in a Product Backlog in the form of tasks. He must manage these tasks by prioritising some of them. He is also responsible for the value and return on investment (**ROI**) of the product. The development team is generally a small one (around seven people [4]). It is capable of self-management. There are no specific roles or positions. Developers have to deal with tasks that are chosen from the Sprint Backlog. Finally, the **Scrum** Master is required to be of service to the team. He accompanies the team of developers and prevents obstacles from getting in the way of progress. It is the **Scrum** Master who establishes the practices and rules of **Scrum**.

2.1.2 Scrum process

The **Scrum** process is relatively simple. It follows the theory described in the official **Scrum** guide [7]. The life cycle is shown in figure 2.3.

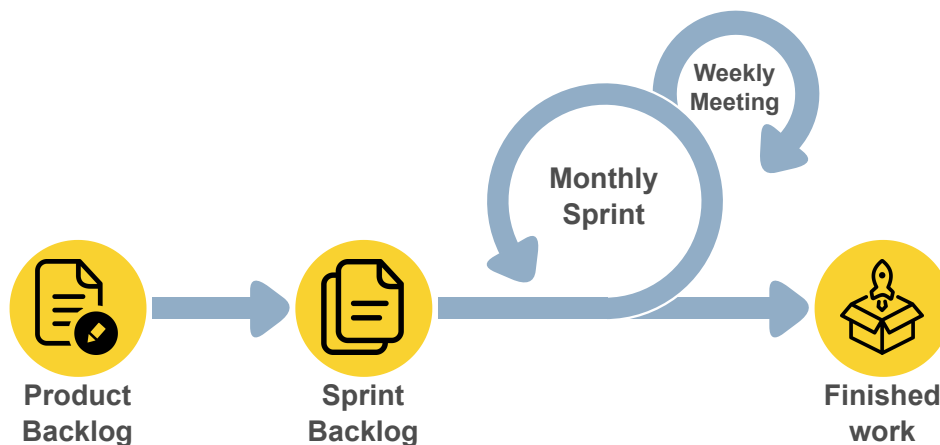


Figure 2.3 *Scrum* life cycle

The process is made up of multiple elements. It begins on the left of the diagram with the Product Backlog. This is a backlog that is filled by the [PO](#). Since the Product Backlog is generally filled with many functionalities divided into tasks, only one part needs to be selected to be placed in the Sprint Backlog. A Sprint generally corresponds to a period of one month during which the tasks in the Sprint Backlog must be completed. During a Sprint, team meetings are organised every day. These are called Daily Meetings and usually last 15 minutes. A Daily Meeting ensures that, at least once a day, the entire team is available to get support on any problems encountered. In this project, only weekly meetings are held, depending on the distance separating the [Scrum](#) Master and the team of developers. Before each Sprint, a Sprint Planning is carried out to determine the tasks from the Product Backlog to be put into the Sprint Backlog. At the end of a Sprint, an inspection is made of how the Sprint went, with the aim of improving quality and efficiency. This is called a Sprint Retrospective. Finally, it's important to remember that at the end of each Sprint, one stage of the final product is completed. Sometimes this is called an increment.

2.2 KanBan methodology

The [KanBan](#) methodology is very well integrated into [Scrum](#). Its aim is to manage the workflow as efficiently as possible. According to the [KanBan](#) guide produced by [Scrum.org](#) [8], several fundamental metrics should be taken into account for the flow:

- **Work In Progress (WIP)** : the number of tasks/items started and not completed.
- **Cycle Time** : the time elapsed between the start and end of a task.
- **Work Item Age** : the time elapsed between the start of the task and now.
- **Throughput** : the number of tasks completed per unit of time.

The average cycle time can be predicted using Little's Law [9]:

$$\text{average cycle time} = \frac{\text{average WIP}}{\text{average throughput}} \quad (2.1)$$

It explains that the more tasks there are in progress, the longer it will take to complete them. A cycle time in **KanBan** can be correlated to a Sprint in the **Scrum** methodology. Using this methodology, it is therefore possible to define the right number of tasks for each Sprint Planning. Note that the more you divide a feature into smaller tasks, the easier it will be to predict the cycle time.

KanBan does more than just limit the **WIP**. It provides a transparent display of the workflow. Tasks can be represented in a table, as shown in figure 2.4. In this case, the table is virtual and can be viewed from anywhere by the whole **Scrum** team.

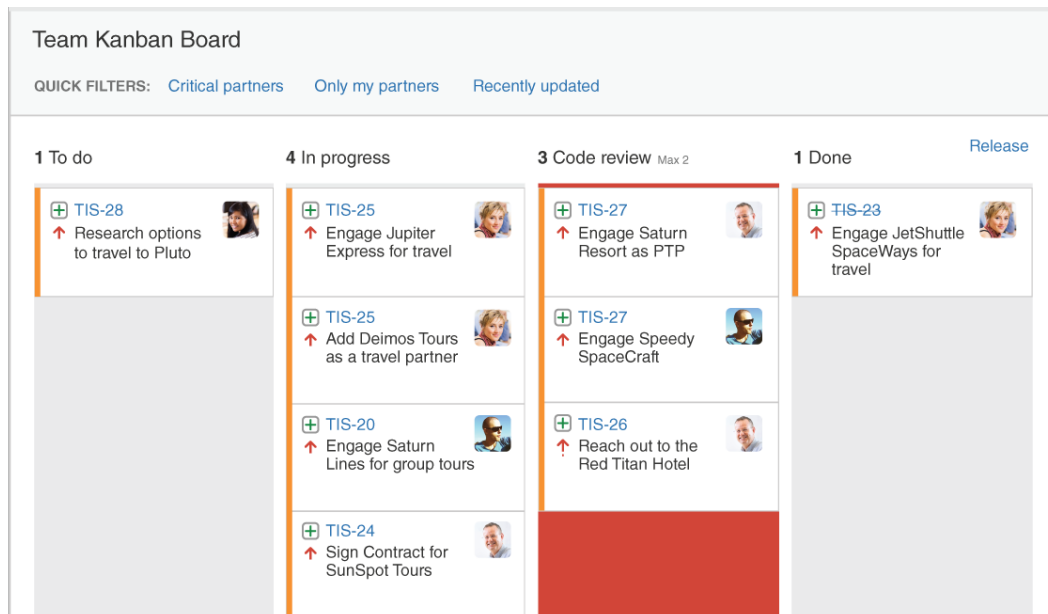


Figure 2.4 Example of a virtual **KanBan** table [10]

Each label in the table corresponds to a task. A task can contain a priority level and is often assigned to a member of the **Scrum** team. All this is defined during Sprint Planning. The labels are categorised in columns to show their progress.

2.3 Project management tools

2.3.1 Version management

Version management, embodied by tools such as Git, remains an essential part of the software development ecosystem. This practice offers a structured approach to tracking and maintaining changes to a project's code. One of the major benefits lies in its ability to keep a complete history of changes, facilitating collaboration between team members working on the same project.




Figure 2.5 Example of version management with two branches

An example of a visual representation of a project is shown in figure 2.5, comprising several distinct versions and branches. Branches allow developers to work on alternative versions of the source code, facilitating parallel development, feature isolation and efficient change management without affecting the main branch.

When part of the development is considered complete, it becomes common practice to create a new version. This version freezes the current state of the project at a specific point in time, providing a clear and stable reference. This ability to create versions means that it is possible to return to previous states of the project if problems arise, whether identified by members of the development team or by external stakeholders, such as customers.

In short, version management, embodied by Git, transcends its simple function of tracking changes to become a central pillar of collaborative working, risk management and preserving the integrity of a software project. Its widespread adoption across the industry is testament to its continuing importance in delivering robust, scalable software development projects.

2.3.2 Continuous Intergration (CI) and Continuous Delivery (CD)

CI/CD (Continuous Intergration/Continuous Delivery) integrates synergistically with version management, and in particular with the DevOps framework, as illustrated in figure 2.6. The aim of the DevOps method is to automate software development processes, thereby encouraging collaboration between development teams (grey area) and operational teams (pink area). Two major components of CI/CD have emerged from this approach, namely Continuous Intergration (CI) and Continuous Delivery (CD), which have become fundamental pillars for developers.




Figure 2.6 DevOps method

CI involves automating various processes with each new software release. These processes typically include building, testing and deploying the software design, as shown in figure 2.6. The main motivation for inventing the CI was to reduce the time needed between launching a project and bringing it to market, a concept known as time-to-market. This automation takes advantage of the constant evolution of machines, contrasting with

the physical and moral limits of human beings. The benefits of **CI** also extend to team collaboration, simplified debugging, reduced costs, improved software quality, and much more.

CD goes hand in hand with the **CI**, aimed at automating the deployment of software to production environments. It eliminates the tedious manual steps associated with deployment, ensuring fast, reliable and frequent delivery of new features and software updates.

CI/CD represents a modern and essential approach to the software development lifecycle, offering efficiency gains, improved team collaboration, accelerated time-to-market and an overall improvement in the quality of software products.

2.3.3 Work packages

Work packages are specific, delimited elements of a project which are assigned to a team or an individual for completion. Each **work package** represents a well-defined task or set of tasks, often associated with concrete deliverables. They facilitate the management and monitoring of the project by allowing a clear division of responsibilities and helping to assess progress. In short, **work packages** are manageable units that contribute to the achievement of an overall project.

2.4 Formations

- AWS Certified Cloud Practitioner

3 | Analysis

The analysis begins by explaining the fundamental concepts that are being worked on throughout the project. Next, a state of the art is provided on the various tools, methods and technologies used in [cloud computing](#) and [IoT](#).

Contents

3.1	Definitions	16
3.1.1	Reference architecture	16
3.1.2	Cloud computing	16
3.1.3	Cloud-Native	17
3.1.4	Internet of Things	18
3.2	Problem formulation	18
3.3	Literature search	19
3.4	History of the Cloud-Native approach	19
3.4.1	The virtualisation	20
3.4.2	The hybrid	20
3.4.3	All to the cloud	20
3.4.4	The Cloud-Native approach	21
3.5	Integrating cloud computing with IoT embedded systems	21
3.5.1	Overview	22
3.5.2	Reference architecture	22
3.5.3	Existing problems and solutions linked to integration	23
3.6	IoT cloud platforms	25
3.7	Cloud infrastructure tools	27
3.7.1	Integrated IaC tools	28
3.7.2	Universal IaC tools	29
3.7.3	Comparison	30
3.8	Arm SystemReady	31
3.8.1	SystemReady certifications	32

3.1 Definitions

Before going into more detail on the subject of this work, it is important to look at the definitions of the various terms frequently used.

3.1.1 Reference architecture

A reference architecture is a solution model in a specific domain. It must be built so that an architecture can be established on its foundations, to make the task of software developers easier.

The aim is to generalise a solution that shows how it works from an overview. It must be possible to observe the relationships and their interactions between the multiple components of the application based on the reference architecture. There are different layers of abstraction depending on the field of application. A high level of abstraction means that it will be possible to understand the solution through more abstract elements such as the general components that will consolidate the application, for example. A low level of abstraction means that the solution will be more precise and certainly more specific to a use case. It will be possible to find detailed relationships between the services contained in the general components.

3.1.2 Cloud computing

The National Institute of Standards and Technology (NIST) [11] has standardised the definition of **cloud computing** as follows:

Cloud computing is a model for enabling ubiquitous, convenient, on-demand network access to a shared pool of configurable computing resources (e.g., networks, servers, storage, applications, and services) that can be rapidly provisioned and released with minimal management effort or service provider interaction. [11]

It is added in the article [11] that this model of **cloud** consists of three service models and four deployment models. The service models are as follows:

- **Infrastructure as a Service (IaaS)** : The capability offered to the consumer consists of providing processing, storage, network and other fundamental computing resources where the consumer can deploy and run arbitrary software, which may include operating systems and applications. The consumer manages the operating systems, storage and applications.
- **Platform as a Service (PaaS)** : The capacity provided to the consumer consists of deploying on the **cloud infrastructure** applications created or acquired by the consumer using programming languages, libraries, services and tools supported by the provider. The consumer manages only the applications and storage.
- **Software as a Service (SaaS)** : The consumer has the option of using the supplier's applications running on a **cloud infrastructure**. The consumer manages nothing of the infrastructure apart from any configuration of the applications.




Figure 3.1 Comparison of different service models [12]

The deployment models are as follows [11] :

- **Private cloud** : The **cloud infrastructure** is made available for exclusive use by a single organisation comprising multiple consumers. It may be owned, managed and operated by the organisation, by a third party or by a combination of both, and it may exist on the organisation's premises or off-site.
- **Community cloud** : The **cloud infrastructure** is reserved for the exclusive use of a specific community of consumers from organisations with common concerns. It may be owned, managed and operated by one or more organisations within the community, by a third party or by a combination of such organisations, and it may exist on or off premises.
- **Public cloud** : The **cloud infrastructure** is made available to the general public for open use. It may be owned, managed and operated by a business, educational institution or government organisation, or a combination of these. It is located on the provider's premises.
- **Hybrid cloud** : The **cloud infrastructure** is a composition of two or more distinct infrastructures (private, community or public) which remain unique entities, but which are linked by a standardised or proprietary technology that allows the portability of data and applications.

3.1.3 Cloud-Native

Cloud-Native is an approach to software development which aims to design, implement and manage applications in the **cloud**. **Cloud computing** environments will be necessary for the proper execution of workloads.

The infrastructure can be deployed in private, public or hybrid [clouds](#). This approach will enable the development of modern, easily scalable applications. The flexibility aspect is also emphasised thanks to techniques for decoupling the multiple services of an application. Another strength of the [cloud](#) is its reliability and sustainability. Since the number of resources is incalculable and they are distributed around the world, redundancy guarantees security in the event of an incident, with almost instantaneous migration of an application's execution. The customer is therefore very often spared any unforeseen events. It is also possible to provide updates in real time and on a recurring basis.

The main benefits are high efficiency, reduced costs and high availability. Applications will be able to exploit resources that are optimal for their use case. The term pay-as-you-go is widely used, as the cost will depend solely on the use made of the resources and not on anything else, such as their maintenance, hardware security, etc. Availability is high because of the astronomical number of resources made available by the various [cloud](#) environment providers.

3.1.4 Internet of Things

[Internet of Things \(IoT\)](#) refers to the interconnection between physical objects and the internet. Objects can range from light bulbs to medical devices and much more. The main areas concerned are home automation and medical technology.

This involves linking different objects and applications to move into an automated world. As this becomes more and more widespread, the services needed to make it work need a huge amount of resources. This has been made possible by the [cloud](#). Everything must be accessible from anywhere in the world, quickly and securely. Several communication technologies are possible to guarantee excellent accessibility and reliability.

3.2 Problem formulation

The general question for analysis is: "Which technologies should be used to develop a reference architecture".

To help answer this question, it is important to consider the following aspects:

- Technology
- Environment
- Tools
- Certification

Since the development uses the [Cloud-Native](#) approach, its history can be described first. The integration of the [IoT](#) world into [cloud computing](#) must then be observed through various scientific works and research. A study of [cloud](#) service providers must be undertaken to differentiate the IoT services offered by each one. As the reference architecture is based on an [Infrastructure as Code \(IaC\)](#), research into [IaC](#) tools is carried out. Finally, it is worth finding out whether SystemReady certification exists for electronic boards using an [Arm](#) processor architecture.

3.3 Literature search

Literature search has focused on academic search engines. The following search engines were used :

- [Google Scholar](#)
- [IEEE Xplore](#)
- [ScienceDirect](#)

The sources were mainly conference papers. Some information was found on websites. The following keywords were used to find sources that met the requirements of this analysis:

- integration
- embedded systems
- [IoT](#)
- [Cloud-Native](#)
- cloud
- [cloud computing](#)
- [cloud providers](#)
- [cloud platforms](#)
- [Infrastructure as Code](#)
- [IaC tools](#)
- [Arm SystemReady](#)

3.4 History of the [Cloud-Native](#) approach

[Cloud-Native](#) is a term that has been around for several years. Figure 3.2 shows the evolution of this term since 2006. The boom took place around 2016. It is undoubtedly due to the birth of Docker (2013) [13] and Kubernetes (2014) [14]. In 2015, the Cloud Native Computing Foundation [15] was created with the aim of making the [Cloud-Native](#) approach ubiquitous [16].




Figure 3.2 Google trends (01.01.2006 until 22.05.2016) of term [Cloud-Native](#) [17]

Before all this, there were only on-site data centres. Generally speaking, each company had its own servers running in its own data centre. This meant that servers were always set up for a specific application. [18]

3.4.1 The virtualisation

The first change came in the 2000s with the virtualisation of servers, although this has been around since the 1960s. When a new application had to be developed, new physical servers had to be bought. With virtualisation, it is no longer necessary to buy new ones. It is possible to run several applications on the same server using virtualisation, which limits the hardware resources for each application (figure 3.3). [18]




Figure 3.3 Traditional server (left) and virtualisation (right)

3.4.2 The hybrid

Despite this, a minimum of one server had to be purchased for an application to work. There were also resource limits for a certain number of applications. Unfortunately, not everyone could afford to buy new servers. In 2006, [AWS](#) launched three web services, [Elastic Compute Cloud \(EC2\)](#), Simple Storage Service (S3) and Simple Queue Service (SQS), to allow organisations and individuals to use Amazon's IT infrastructure on an as-needed basis at basic prices [19]. This initiative stems from a restructuring of the Amazon platform for better scalability and this has made it possible to sell virtual servers as a service [20]. By extension, the myth of making money from servers not in use for the majority of the year has been debunked by Benjamin Black, co-founder of the [EC2](#) [20]. This is where the world of [cloud](#) really began. Kratzke and Quint confirm this in their scientific article [17]. The purpose of the [EC2](#) service, which is still used today, is to offer a virtual machine in the [cloud](#). This is a [cloud computing](#) environment for running workloads. It is now possible to migrate your infrastructure from the database to the [cloud](#). The corresponding expression is "Lift and Shift" [21]. These were undoubtedly the first so-called [Cloud-Native](#) approaches [18]. However, this term only appeared in papers for the first time in 2012 [17]. These were two conference papers proposing solutions for [Cloud-Native](#) applications [22, 23]. Organisations have learned a lot about only paying for what you use. From there, there's a move to hybrid where companies are still using their on-premises servers alongside servers in the [cloud](#). [18]

3.4.3 All to the [cloud](#)

A new barrier has now been crossed. It's no longer a question of doing hybrid, but of transferring the entire infrastructure to the [cloud](#). This was made possible when [cloud](#) service providers such as [AWS](#), Microsoft Azure, Google Cloud and many others developed several services for all types of use. The reason there are different services is quite simply to decouple application functionality as much as possible. This is known

as a microservices architecture. The risk of an entire application being interrupted in the event of a problem is much less likely. However, some organisations still have applications that are several decades old and have a monolithic architecture. It is not necessarily possible to decouple functionalities. That's why they work on a hybrid basis. [18]

3.4.4 The Cloud-Native approach

If we ask several engineers today about the definition of the Cloud-Native approach, we will find varying interpretations. There are currently three main schools of thought on this subject [18]. The first group considers that an approach is Cloud-Native when all the workloads run in the cloud. A minority within this group argue that you can be partially Cloud-Native, i.e. have one complete application deployed in the cloud while maintaining another application locally. However, others feel that this is still a hybrid approach. A second group says that to be truly Cloud-Native, you need to fully exploit the capabilities offered by the cloud. This means not only using the basic services of a cloud provider, but also making full use of the advanced features on offer, such as serverless functions. Some even consider that a company has to be born in the cloud to be truly Cloud-Native, a phenomenon that is increasingly being observed, with companies launching their first applications directly in the cloud without ever having used on-premises servers. This is the case for the company 56K.Cloud To be confirmed. [18]

In 2018, a brief history of cloud application architectures described the term Cloud-Native as follows :

Cloud infrastructures (IaaS) and platforms (PaaS) are built to be elastic. Elasticity is understood as the degree to which a system adapts to workload changes by provisioning and de-provisioning resources automatically. Without this, cloud computing is very often not reasonable from an economic point of view. Over time, system engineers learned to understand this elasticity options of modern cloud environments better. Eventually, systems were designed for such elastic cloud infrastructures, which increased the utilization rates of underlying computing infrastructures via new deployment and design approaches like containers, microservices or serverless architectures. This design intention is often expressed using the term Cloud-Native. [24]

3.5 Integrating cloud computing with IoT embedded systems

With technology advancing at a rapid pace, today's world appreciates help in making everyday life better [25]. Embedded systems that have become increasingly connected now offer this service. However, since they need to attend any location for any domain and task, they need to be small for better integration. As a result, resources are limited for computing operations, data storage, data processing and so on. The built-in security and low energy consumption that these devices must also ensure are not enough to allow them to move forward. This is why cloud computing needs to be integrated to fill the gaps left by IoT [25]. Cloud of Things is one of the paradigms given to the combination of these two technologies in 2014 [26].

3.5.1 Overview

A model for integrating **cloud computing** into an embedded system is shown in figure 3.4. [27]

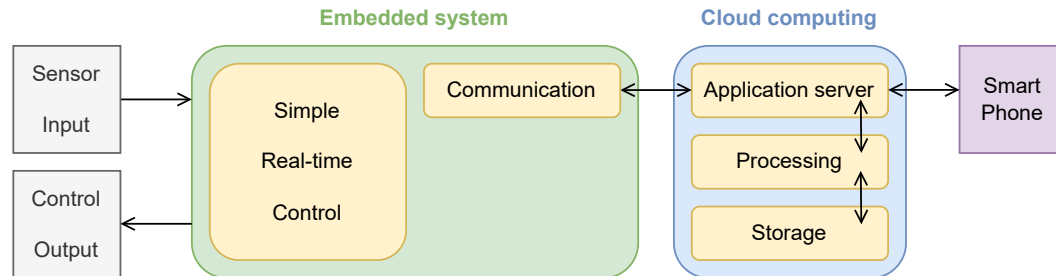


Figure 3.4 Model for integrating **cloud computing** into an embedded system [27]

The main idea is to use only what is necessary on the **IoT** device. In other words, all computing and storage functions should be moved to the **cloud**. The device should only be linked to the input and output peripherals using a simple controller and establish communication with the **cloud**. For the integration to work, the embedded system must have an internet connection. **Cloud computing** would contain the core of the application as well as various services to perform analysis, processing and calculation operations, and store data in real time. From there, it would also be possible to view the data from a smartphone. This approach is a solution that was thought up by Furuichi and Yamada in 2014 [28]. It offers a number of advantages, such as reduced energy consumption by eliminating large workloads and reduced size by freeing up electronic components. The **cloud** also makes it easy to scale up if necessary. The application can then be highly scalable. [27]

Furuichi and Yamada wanted to use a project to prove that their suggested approach worked. They developed a traffic jam detection system that recognises car licence plate numbers along with their location and time. For a plate number to be recognised from a raw image, image processing is carried out using machine learning. They found that this application, split between an **IoT** device and a **cloud computing** environment, scored much better than the same application managed entirely on a laptop or embedded system. The evaluation looked at cost, battery life, performance, scalability and reliability. [28]

3.5.2 Reference architecture

A reference architecture has been developed to support **IoT** objects in **cloud computing** (figure 3.5). The Standards Development Organization has decided to make this a standard by launching the Cloud Standards Customer Council programme to drive forward the adoption of **cloud computing**. In this architecture, various aspects are taken into account: scalability, security, reliability and protection of privacy. [29]

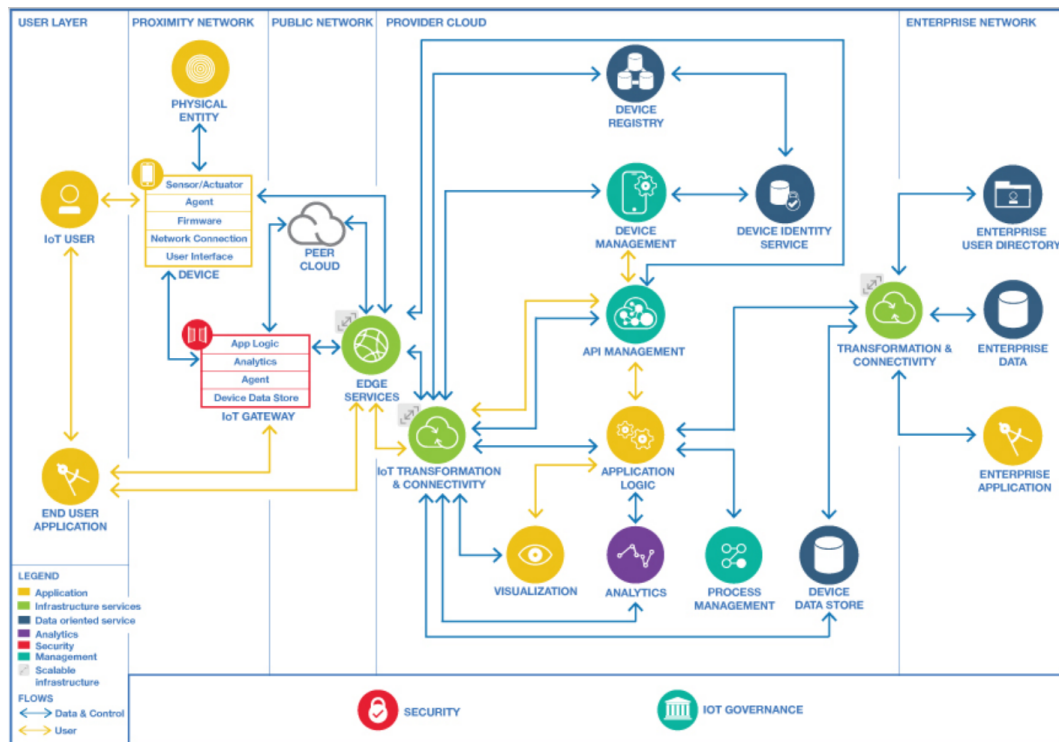


Figure 3.5 Cloud Customer Reference Architecture for IoT [29]

The architecture proposed in figure 3.5 is made up of components and their relationships. They are separated under a three-tier architecture model, called 3-Tier (presentation, logic, data). There is edge-tier, the platform tier and the enterprise tier.

The edge-tier contains the "Proximity Network" and "Public Network" parts. It represents the place where data is collected. It is collected from an IoT device in a proximity network. The data will either be sent to an IoT gateway or directly to the cloud provider.

The cloud provider tier manages the entire logical part of the application. It includes the collection, processing and analysis of data flows. Each device is also managed in this tier, as is its identity, for example.

The enterprise tier contains the "Enterprise Network" part. It includes the enterprise data and the application. Data from cloud computing can be stored in enterprise data.

3.5.3 Existing problems and solutions linked to integration

Integrating IoT and cloud computing poses a number of challenges. They are to be found both on the hardware side and in the cloud. [26]

There are many communication protocols. There are protocols for every type of use. On the IoT side, for example, there are Zigbee [30], Bluetooth and Bluetooth Low Energy (BLE) [31], LoRaWAN (long range) [32], WiFi, Thread [33] and the latest Matter protocol [34]. In an IoT network, there can be a multitude of sensors linked to a gateway. Interoperability is a problem in this case. Since 2022, version 1.0 of the Matter protocol has been published to address this problem. However, few devices are yet compatible. In research projects, Zigbee is very popular [25, 26, 35, 36]. Two main technologies are used

for communication between the **IoT** gateway and the **cloud**: **Message Queuing Telemetry Transport (MQTT)** [37] and **Constrained Application Protocol (CoAP)** [38]. Several projects have used **CoAP** for resource reasons [39, 40, 41]. It is a protocol dedicated to constrained devices with small amounts of memory [38], using **UDP**. It is also suitable for devices that need to bind to web services as it uses the **REpresentational State Transfer (REST)** model. **REST** is an architectural style for web services [39]. It would have been possible to use the **Service-Oriented Architecture (SOA)** style, but it is less suited to this use case [39]. Service-based architecture styles allow a high degree of decoupling from the application. This is an advantage in the event of a problem with a service, so as not to completely freeze the application. Finally, **MQTT** is specially designed for **IoT** [37]. It uses the **TCP** transport protocol. **Cloud** provider **AWS** adopts the **MQTT** protocol in its **IoT** sector [42]. A communication between an **Arm** microcontroller and **AWS** proved that this protocol works [41]. In fact, the data passing through is often in **JavaScript Object Notation (JSON)** format [27, 36, 43]. This is a lightweight data representation syntax for storing and exchanging textual information [36]. There is also the **HyperText Transfer Protocol (HTTP)** communication protocol. It would be suitable for updating **IoT** devices. It allows more bytes to be sent per packet, which increases the update speed. Furthermore, in industry, **HTTP** is generally less constrained by firewalls than **MQTT** or **CoAP** [44].

A lot of energy is consumed in **IoT** devices. The reason for this is the increasing amount of data being transferred to the **cloud**. Energy needs to be managed efficiently, for example by setting a sleep mode. If this is not possible, natural energy could be used to provide power, such as solar, wind or vibration. [26]

Resource allocation is another debate. Each **IoT** object may have a different function. They do not necessarily require the same resources in the **cloud**. What's more, unexpected events could occur requiring a greater workload. The requirements could not necessarily be predictable. Even if they were, the cost of increasing unused resources for most of the time would be a loss. **Cloud** service providers offer a number of solutions to this problem. For example, **AWS** offers a service that balances the load in the **cloud computing** environment [45]. The combinatorial auction approach is popular for resource allocation in the **cloud** [46]. It must satisfy **Quality of Service (QoS)** constraints and maximise the **cloud** provider's profit. Provider and user profit, resource utilisation and quality of service are the dominant performance factors. Research has succeeded in optimising this approach [46]. Nevertheless, as soon as a new **IoT** node is added, a sample of data could be sent to the **cloud** to set up the minimum resources required. [26]

Identity management in an internet network is very important to ensure that each device is unique. Given the number of **IoT** devices today, **IPv6** addressing would be the optimal solution. It would be sufficient to support this type of network. However, deployment remains problematic. The coexistence of **IPv6** with **IPv4** is not so easy [26]. There are three possible techniques for cooperation between the two [47]. The **IoT** device may support both types of addressing. Alternatively, **IPv6** packets must be encapsulated in **IPv4** packets, or a **Network Address Translation (NAT)** must be available that translates both types of packet.

The discovery of new IoT devices to be integrated into the cloud needs to be managed. At any time, a device can be part of a service and leave it. It is also necessary to monitor its status once it is connected and keep it up to date. AWS enables a component to be deployed on IoT devices that links to the cloud. From there, it is possible to perform Over-The-Air (OTA) updates [48] and monitor the state of the device [49]. Azure also provides these different services [50, 51]. An update method was designed in this area in 2022 [44]. It is based on the standard update architecture for IoT [52], proposed by the Internet Engineering Task Force (IETF).

Cloud computing environments must respond correctly to data flows, whether large or small. The urgency with which certain data is transmitted is also an integral part of this QoS. QoS is assessed in terms of bandwidth, delay, packet loss rate and delay in the transmission of data packets (jitter). [26]

Another point to consider is where the data is stored. Depending on the volume of data, it is best to store it in the physical location closest to the user. Sensitive data must also be managed in accordance with each country's data protection laws. Fortunately, cloud providers offer the option of choosing the region where data is stored. [26]

Security seems to be an issue in integration. In 2016, there was a lack of trust in cloud service providers and a lack of awareness of Service-Level Agreements (SLAs) [53]. A Service-Level Agreement is a document defining the provision of cloud services and the responsibilities of the provider and the customer. 56K.Cloud claims some negligence with developers. However, when data is transferred to the cloud, the location of databases is not always transparent. Distributing data over several locations, while ensuring high availability, also increases the chances of sensitive data being leaked. A distributed system of this kind is exposed to various potential attacks, such as SQL injections, cross-site scripting attacks and many others. Significant vulnerabilities can also be exposed, including session hijacking and virtual machine evasion. In addition, the computing power constraints associated with connected objects limit the application of public key cryptography to all layers of the system. Two years later, in 2018, in an attempt to partially address this problem, two security models have been proposed using two encryption algorithms (AES and RSA) [35]. They can be used to integrate IoT and cloud computing. However, they remain sub-optimal and could be the subject of future research.

Unnecessary data communication is a current trend. IoT devices generate all kinds of data. It would be interesting to have an IoT gateway that manages traffic by letting data circulate only when it is necessary and only that which is necessary [26]. This would save energy and make better use of network and cloud resources. An architecture for this approach has been presented [54].

3.6 IoT cloud platforms

A number of projects have created their own cloud platform [55]. The University of Glasgow decided to create its own cloud computing environment based on a Raspberry Pi [56], called PiCloud [57]. It emulates every layer of a cloud stack, from resource virtualisation to network behaviour. A fleet of 56 Raspberry Pi devices was able to consolidate this cloud platform. This project is ideal for education, thanks to its low

cost. Another study, based on the previous one, carried out the same style of [cloud](#) platform, but with 300 Raspberry Pi devices. Called the Bolzano Raspberry Pi [58], this [cloud computing](#) environment incorporates several [Network Attached Storage \(NAS\)](#) as storage units. However, these are traditional [cloud infrastructures](#).

[Cloud infrastructures](#) have been adapted for the use of [IoT](#). What differentiates them from traditional infrastructures is the processing of data generated by events in real time [59]. There are now a large number of public [cloud](#) service providers for [IoT](#). There are also a number of open source [IoT](#) solution providers. The main ones today are [AWS](#), Microsoft Azure and Google Cloud [60].

[AWS](#) [61] provides a [cloud](#) platform with over 200 services. Resources are distributed across 32 regions worldwide. It claims to be the most widely adopted platform in the world. Products include: compute, storage, databases, analytics, networking, mobile, development tools, management tools, [IoT](#), security and enterprise applications. It offers its services on demand and payment is on a pay-as-you-go basis. Through a survey published in 2019 on the various platforms [62], a solution dedicated to [IoT](#), called [AWS IoT](#), was summarised. This service aims to collect, store and analyse data from devices. Amazon FreeRTOS and [AWS IoT Greengrass](#) are made available by [AWS IoT](#), facilitating the development of embedded applications on these devices. For embedded systems, both [Arm](#) and x86 architectures with Linux are supported [63]. For device management, there is [AWS IoT Core](#), which facilitates device connectivity with [cloud](#) services. This can be complemented by [AWS IoT Things Graph](#) to visualise the data emitted by devices. In terms of data processing, the platform provides [AWS IoT Analytics](#), offering developers a convenient way to analyse the data generated. The results of this process can be routed to other devices or systems via the [AWS IoT Events](#) service [62] service. The aim of a thesis project was to define an [IoT](#) architecture for connecting [Arm](#) microcontrollers to [AWS](#) [41]. An example of a smart home application was made on this architecture. A machine learning service was used in the [cloud computing](#) to make predictions from data sent from the [IoT](#) devices. The company 56K.Cloud is also carrying out a few projects linking [Arm](#) devices with a [cloud infrastructure AWS](#). Typically, it uses the [AWS IoT Greengrass Core](#) component to run applications on embedded systems.

Microsoft Azure [64] is another public [cloud](#) provider. Launched in 2008, it offers more than 200 [cloud](#) products and services. Its data centres are located in more than 60 regions around the world. The company operates in the healthcare, finance, public sector, manufacturing and retail sectors. It says it invests a billion dollars a year in security to protect customer data. As far as [IoT](#) is concerned, it offers a solution called [Azure IoT](#). This solution includes several services. [Azure IoT Hub](#) is a service that establishes connections, administers and develops the ability to manage billions of [IoT](#) devices, from the peripheral to the [cloud](#). [Azure IoT Central](#) provides a user interface and [APIs](#) for connecting and managing large-scale embedded systems. Data is visualised using [Azure Time Series](#). [Azure Sphere](#) offers security to protect [IoT](#) devices. Physical device spaces can be replicated using the [Azure Digital Twins](#) service. System logic can be managed by [Azure IoT Edge](#). For the development of integrated applications, there is [Azure RTOS](#). [64]

Google Cloud [65] is one of the world's leading cloud providers. It currently distributes its resources in 39 regions across the globe. Launched in 2008, the provider is trying to keep up with the competition by offering around a hundred services for all types of business. However, Google has announced that it will be pulling out of the IoT business in 2023. The Google Cloud spokesperson said that the needs of their customers could be better served by specialist IoT partners [66].

There are other cloud providers offering IoT services. Some of these are mentioned in figure 3.6, where a comparison is made.

IoT platform	Open source	Protocols	Data store	Push notification	Trigger	Visualisation
AWS IoT	No	MQTT, HTTPS	Yes	Yes	Yes	Yes
Azure IoT	No	MQTT, AMQP, HTTPS	Yes	Yes	Yes	Yes
IBM Watson IoT Platform	No	MQTT, JMS	Yes	Yes	Yes	Yes
Kaa Enterprise IoT Platform	Yes	MQTT, HTTPS	Yes	Yes	Yes	Yes
ThingsBoard	Yes	MQTT, CoAP, HTTP, LwM2M, SNMP	Yes	Yes	Yes	Yes

Figure 3.6 Comparison of IoT cloud platform providers (sources from providers)

Looking at the back end of platforms, they would have an advantage in using containerisation over traditional virtual machines [67]. It is seen as a lightweight virtualisation solution, with greater flexibility. In addition, containers are proving particularly suitable for solving the platform issues commonly associated with PaaS services in the cloud. This includes aspects such as application packaging and coordination [67]. Docker is the most popular container solution [13]. As for container orchestration, there's Kubernetes, an open source system [14].

3.7 Cloud infrastructure tools

When you want to host applications in the cloud, you need to deploy a cloud infrastructure. It is perfectly possible to build the infrastructure manually. All you have to do is visit a cloud provider's web page and select the resources you need from the services on offer. However, this infrastructure could be deployed several times. Fortunately, there are tools that can automatically deploy cloud infrastructures, called Infrastructure as Code (IaC). The automatic deployment of cloud infrastructures corresponds to provisioning. These tools use templates containing a description of the infrastructure based on executable code or a configuration file. They save time on deployment and ensure that the implementation policy behind the infrastructure is always the same (security, rules, etc.). In addition to provisioning, IaC tools are capable of updating and deleting an infrastructure. There are universal provisioning tools for different cloud providers and there are integrated tools for each provider. [68]

3.7.1 Integrated IaC tools

Having mentioned the [AWS](#) and Microsoft Azure providers in section 3.6, each of them has its own [IaC](#) tool reserved for use with its platform. Google [Cloud](#) is not mentioned in this section due to the cessation of [IoT](#) activities.

At [AWS](#), there are two [IaC](#) tools. [AWS](#) CloudFormation [69] is the first. [AWS Cloud Development Kit \(CDK\)](#) [70] is the second for defining and building infrastructure in the [AWS cloud](#) environment. [AWS CDK](#) gives users the flexibility to define infrastructure in distinct programming languages in the form of an application, with imperative syntax. It uses [AWS](#) CloudFormation as the deployment engine and allows users to define the infrastructure using programming idioms to model the system design. The [CDK](#) application deployment process includes the construction of the defined elements, preparation, validation and synthesis of the deployment artefacts. The [CDK](#) application is then transferred to the [AWS](#) CloudFormation service for real deployment. An overview of the deployment is shown in figure 3.7. The basic elements of [CDK](#) applications are called "constructs", and they represent [cloud](#) components containing services. Constructs can be nested to create a hierarchy of dependencies called a "construct tree". Ultimately, this hierarchy defines how constructs are synthesized into resources in the final [AWS](#) CloudFormation model. In a [CDK](#) application, it is possible to define several stacks. These are unique deployment units, each containing a hierarchy of constructs. [68]

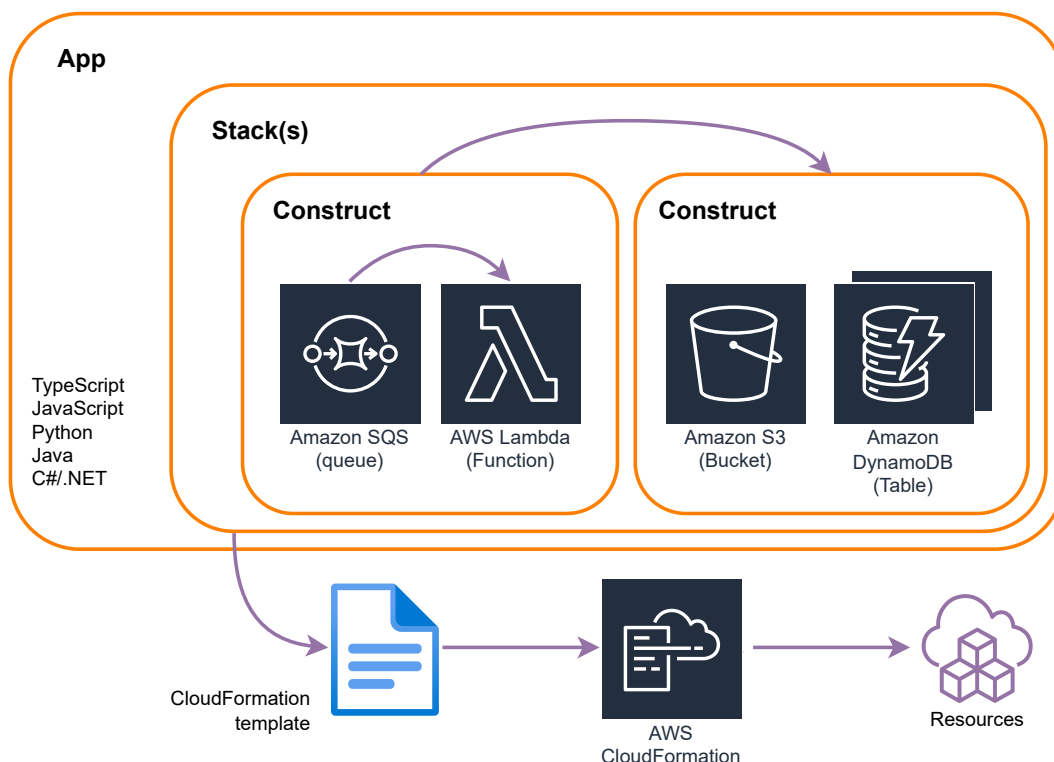


Figure 3.7 Overview of [AWS CDK](#) [70] deployment [70]

Microsoft Azure offers the **Azure Resource Manager (ARM)** deployment and management service [71]. All actions performed on resources by the user go through the ARM manager. It authenticates and authorises them. At the deployment level, there are ARM models, which are JSON files defining the infrastructure and configuration of a cloud environment. JSON uses declarative syntax. Before deployment, the ARM API checks the models for possible errors, guaranteeing successful deployment. ARM models ensure an identical architecture for deployments in different environments, such as development, testing and production, for example. The fewer dependencies there are, the faster ARM can deploy infrastructure resources in parallel. ARM models can be segmented into modular files for easy reuse. A group of resources deployed in the Azure platform can be extracted in the form of an ARM model. [68]

3.7.2 Universal IaC tools

There are several IaC tools compatible with different cloud platforms such as AWS, Microsoft Azure, and many others. The 56K.Cloud company mainly uses Terraform [72] and Pulumi [73].

Terraform is a tool created in 2014 by HashiCorp, written in the Go programming language [74]. It mainly provisions **Infrastructure as a Service**, but can also deploy **Platform as a Service** and **Software as a Service**. Terraform works using two inputs: configuration and state. The desired resources in the cloud infrastructure are described in a configuration file. This file works with the Terraform language using declarative syntax. The state corresponds to the current state of the infrastructure, or in other words the stack. The state is managed by the Terraform API and is stored locally by default. The deployment engine compares the desired infrastructure with the current state and determines which resources need to be created, updated or deleted. Terraform uses provider plugins to interact with cloud providers. It offers five basic commands for different stages: "init", "validate", "plan", "apply" and "destroy", enabling the infrastructure to be managed efficiently and easily [72]. The "init" command prepares the project directory. The "validate" command checks that the configuration is valid. The "plan" command displays the changes required by the current configuration. The "apply" command creates or updates the cloud infrastructure. Finally, the "destroy" command destroys the cloud infrastructure. An overview of the deployment can be seen in figure 3.8. [68]

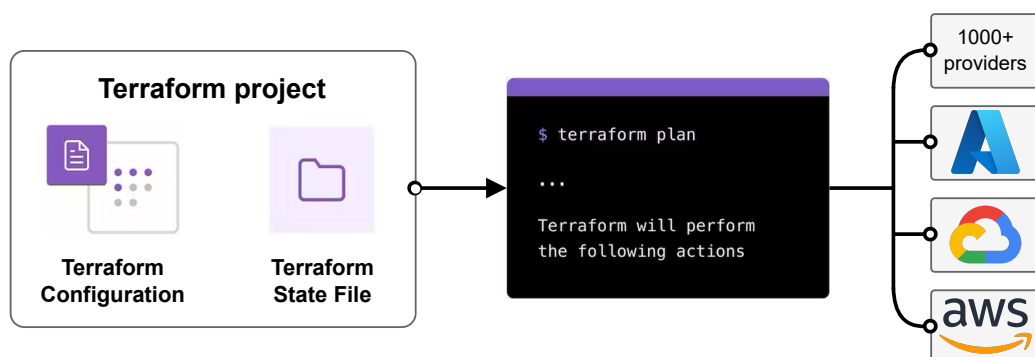


Figure 3.8 Overview of Terraform deployment [72]

Pulumi is an [IaC](#) tool similar to Terraform, but open source. It is written in the Go language [75]. Version 1.0 of Pulumi was released in 2019. This tool can be used with several programming languages, such as TypeScript, JavaScript, Python, Go and .NET. It therefore offers developers the possibility of creating [cloud infrastructures](#) using standard technologies, unlike Terraform. Deployment follows the Terraform model, where a Pulumi application is run to schedule the desired resources in the [cloud infrastructure](#). The deployment engine compares these desired resources with the current state of the stack and determines the necessary actions. The state of the last infrastructure deployment is saved in the Pulumi [Cloud](#) platform by default. With the Pulumi tool, we're talking about applications, because it uses imperative syntax. Applications describe how the [cloud infrastructure](#) should be composed by allocating resource objects whose properties correspond to the desired state. Resource operations are executed in parallel if possible, but some resources have dependencies between them. Deployment is managed by the Pulumi [CLI](#). An overview of deployment is shown in figure 3.9. [68]

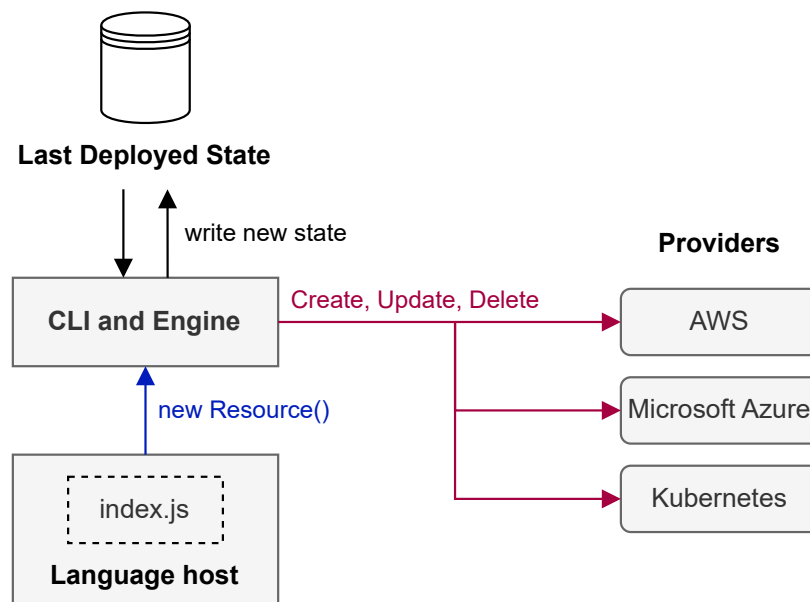


Figure 3.9 Overview of Pulumi deployment [73]

3.7.3 Comparison

Without even using the [IaC](#) tools, it is possible to identify a few differences. The tools integrated with the [cloud](#) providers cannot be used outside their platform. The other two tools, Terraform and Pulumi, can be provisioned by different providers. One notable point concerns the syntax. Terraform and Microsoft Azure use descriptive syntaxes. These limit the variety of programming languages. Terraform has its own language which you need to familiarise yourself with the first time. [AWS](#) and Pulumi use imperative syntax, which leaves the choice of programming languages open.

Focusing on the practical side, other disparities were noted. These stem from a project that was carried out with the aim of comparing the different [IaC](#) tools mentioned in this section 3.7 [68]. Terraform has several advantages, including clear documentation, an understandable configuration language, real-time execution plans and the ability to structure the infrastructure by dividing scripts into separate files. Because [cloud](#) providers

have different services, it is impossible to reuse the same code with each of them. It is possible to maintain the overall architecture of the infrastructure if the entities are clearly separated in the scripts. On the Pulumi side, resources can be separated into classes or functions. Common code between providers is therefore limited to abstractions of classes and functions. Pulumi has developed a tool capable of converting Terraform scripts into Pulumi scripts. It differs from Terraform by using Pulumi [Cloud](#) for infrastructure state management by default instead of a self-managed [API](#). This is an advantage of not needing to store state files locally and manage them. Terraform and Pulumi generally use [APIs](#) from [cloud](#) providers to create resources. One case stands out with Pulumi, when it works with [AWS](#), it uses the [AWS SDK](#) service directly, which will create the resources. Concerning the [Azure Resource Manager](#) tool, it is less efficient than [AWS CDK](#) because of the complexity of the configuration files. [AWS CDK](#) offers better readability by being able to divide up the code as required. Despite all this, it should be noted that these four [IaC](#) tools were able to provision the project in question. [68]

3.8 Arm SystemReady

Manufacturer [Arm](#) has decided to introduce a certification programme called SystemReady [76]. In 2020, as part of [Arm's](#) Cassini project, SystemReady was introduced to address the compatibility concerns of users planning to move from x86 architectures to systems based on [Arm's](#) processors [77]. [Arm](#) defines its programme as follows:

Arm SystemReady is a compliance certification program based on a set of hardware and firmware standards: [Base System Architecture \(BSA\)](#) and [Base Boot Requirements \(BBR\)](#) specifications, plus a selection of supplements. This ensures that subsequent layers of software also 'just work'. The compliance certification program tests and certifies that systems meet the SystemReady standards, giving confidence that operating systems [OS](#) and subsequent layers of software just work. [76]

The [Arm BSA](#) specification establishes a hardware basis for system software, including [OSs](#), hypervisors and firmware, based on the 64-bit [Arm](#) architecture. This specification ensures that users can install, boot and run generic [OSs](#) and hypervisors [78]. The [BBR](#) specification complements [BSA](#) by defining the basic firmware requirements necessary to support any [OS](#) or hypervisor compatible with the [BSA](#) specification [79].

A number of advantages were mentioned in relation to the program [76]:

- Compliance allows [OSs](#) and workloads to operate seamlessly on different [Arm](#) platforms.
- Standardisation of a range of different devices and systems provides a stable basis and a choice of systems for all sectors.
- Compliance with standards inspires confidence in software compatibility, so developers can concentrate on innovation and adding value to their products.
- Compliance simplifies and accelerates time-to-market for partners.
- Compliance makes it easy to identify [Arm](#)-based devices thanks to the "System-Ready certified" stamp.

- Compliance makes it easy to deploy and maintain standard firmware interfaces, reducing maintenance costs.
- Compliance reduces the costs associated with adopting a new platform by eliminating custom firmware engineering.

3.8.1 SystemReady certifications

To be able to adapt to different types of devices and markets, [Arm](#) has introduced four certifications. The first is SystemReady SR. It ensures the smooth operation of servers or workstations on a [Arm](#) chip. It is specially designed for Windows, Linux, VMware and [Berkeley Software Distribution \(BSD\)](#) environments. It targets generic out-of-the-box OSs and also supports legacy OSs on new devices. [76]

SystemReady ES is designed to meet the needs of Windows, Linux, VMware and [BSD](#) ecosystems based on [Arm](#) embedded systems. It also targets generic out-of-the-box OSs and supports legacy OSs on new devices. [76]

SystemReady IR ensures that Linux and [BSD](#) work perfectly on [Arm](#) embedded systems. It is ideally suited to the [IoT](#) sector. It mainly targets the Linux environment, but also custom images (Yocto, OpenWRT, buildroot) and pre-built images (Debian, Fedora, SUSE). [76]

SystemReady LS specifies the correct operation of Linux OSs on [Arm](#) chips designed for servers. This certificate is mainly for hyperscalers. Hyperscalers are large [cloud](#) service providers. [76]

At present, a multitude of hardware design companies are partnering with [Arm](#) to follow the SystemReady standards. Several embedded systems are already certified and on the market.

4 | Design

The design provides a theoretical explanation of the implementation of the ...

Contents

4.1	Section 1	34
-----	-----------	----

4.1 Section 1

5 | Implementation

The implementation explains ...

Contents

5.1	Section 1	36
-----	-----------	----

5.1 Section 1

6 | Validation

This part presents ...

Contents

6.1	Section 1	38
-----	-----------	----

6.1 Section 1

7 | Conclusions

- 7.1 Project summary**
- 7.2 Comparison with the initial objectives**
- 7.3 Encountered difficulties**
- 7.4 Future perspectives**

A | Plannings

A.1 Forward planning

[illegible]

A.2 Effective planning

Bibliography

- [1] *56K.Cloud*. Let's Start your Cloud Journey. URL: <https://www.56k.cloud> (visited on 09/21/2023).
- [2] *Freepik*. Images by macrovector. URL: <https://www.freepik.com/author/macrovector> (visited on 09/29/2023).
- [3] Noura Abbas, Andrew M. Gravell, and Gary B. Wills. "Historical Roots of Agile Methods: Where Did "Agile Thinking" Come From?" en. In: 2008. ISBN: 978-3-540-68255-4. DOI: [10.1007/978-3-540-68255-4_10](https://doi.org/10.1007/978-3-540-68255-4_10).
- [4] Elena Mugelini and Omar Abou Khaled. *MSE university course - IT Project Management*. PHASE III: execute the IT project - choosing a SDLC methodologies. 2022.
- [5] *Manifesto for Agile Software Development*. URL: <https://agilemanifesto.org/iso/en/manifesto.html> (visited on 10/05/2023).
- [6] *Scrum Guides*. Jeff Sutherland. URL: <https://scrumguides.org/jeff.html> (visited on 10/05/2023).
- [7] Ken Schwaber and Jeff Sutherland. *The Scrum Guide*. 6th ed. 2020.
- [8] Scrum.org, Daniel Vacanti, and Yuval Yeret. *The Kanban Guide for Scrum Teams*. 2021.
- [9] *Wikipedia - Little's law*. URL: https://en.wikipedia.org/wiki/Little%27s_law (visited on 10/06/2023).
- [10] *Atlassian - Kanban*. URL: <https://www.atlassian.com/agile/kanban> (visited on 10/06/2023).
- [11] Peter Mell and Timothy Grance. "The NIST Definition of Cloud Computing". en. In: 2011. DOI: [10.6028/NIST.SP.800-145](https://doi.org/10.6028/NIST.SP.800-145).
- [12] *RedHat - IaaS vs. PaaS vs. SaaS*. URL: <https://www.redhat.com/en/topics/cloud-computing/iaas-vs-paas-vs-saas> (visited on 10/13/2023).
- [13] *Docker*. URL: <https://www.docker.com> (visited on 10/12/2023).
- [14] *Kubernetes*. URL: <https://kubernetes.io> (visited on 10/12/2023).
- [15] *Cloud Native Computing Foundation*. MAKE CLOUD NATIVE UBIQUITOUS. URL: <https://www.cncf.io> (visited on 10/12/2023).
- [16] *Cloud Native Computing Foundation*. CNCF Charter. URL: <https://github.com/cncf/foundation/blob/main/charter.md> (visited on 10/12/2023).
- [17] Nane Kratzke and Peter-Christian Quint. "Understanding cloud-native applications after 10 years of cloud computing - A systematic mapping study". en. In: 2017. DOI: [10.1016/j.jss.2017.01.001](https://doi.org/10.1016/j.jss.2017.01.001).
- [18] *Rancher - The History of Cloud Native*. URL: https://www.suse.com/c/rancher_blog/the-history-of-cloud-native/ (visited on 10/12/2023).

Bibliography

- [19] Simson Garfinkel. "An Evaluation of Amazon's Grid Computing Services: EC2, S3 and SQS". en. In: 2007.
- [20] Benjamin Black — *EC2 Origins*. URL: <http://blog.b3k.us/2009/01/25/ec2-origins.html> (visited on 10/31/2023).
- [21] IBM. What is lift and shift? URL: <https://www.ibm.com/topics/lift-and-shift> (visited on 10/12/2023).
- [22] Vasilios Andrikopoulos, Christoph Fehling, and Fran Leymann. "DESIGNING FOR CAP - The Effect of Design Decisions on the CAP Properties of Cloud-native Applications". In: *Proceedings of the 2nd International Conference on Cloud Computing and Services Science - CLOSER*. INSTICC. SciTePress, 2012, pp. 365–374. ISBN: 978-989-8565-05-1. DOI: [10.5220/0003931503650374](https://doi.org/10.5220/0003931503650374).
- [23] Sergio García-Gómez et al. "4CaaS: Comprehensive Management of Cloud Services through a PaaS". In: 2012, pp. 494–499. DOI: [10.1109/ISPA.2012.72](https://doi.org/10.1109/ISPA.2012.72).
- [24] Nane Kratzke. "A Brief History of Cloud Application Architectures". en. In: *Applied Sciences*, 2018, p. 1368. DOI: [10.3390/app8081368](https://doi.org/10.3390/app8081368).
- [25] Manuel Díaz, Cristian Martín, and Bartolomé Rubio. "State-of-the-art, challenges, and open issues in the integration of Internet of things and cloud computing". en. In: *Journal of Network and Computer Applications*, 2016, pp. 99–117. DOI: [10.1016/j.jnca.2016.01.010](https://doi.org/10.1016/j.jnca.2016.01.010).
- [26] Mohammad Aazam et al. "Cloud of Things: Integrating Internet of Things and cloud computing and the issues involved". en. In: 2014, pp. 414–419. DOI: [10.1109/IBCAST.2014.6778179](https://doi.org/10.1109/IBCAST.2014.6778179).
- [27] Nakita Oza. "Integration of Cloud with Embedded Systems for IoT devices". en. In: *International Journal of Latest Transactions in Engineering and Science*, 2018.
- [28] Tetsuo Furuichi and Kunihiro Yamada. "Next Generation of Embedded System on Cloud Computing". en. In: *Procedia Computer Science*, 2014, pp. 1605–1614. DOI: [10.1016/j.procs.2014.08.244](https://doi.org/10.1016/j.procs.2014.08.244).
- [29] Cloud Standards Customer Council. "Cloud Customer Architecture for IoT". en. In: (2016).
- [30] Zigbee. The Full-Stack Solution for All Smart Devices. URL: <https://csa-iot.org/all-solutions/zigbee/> (visited on 10/20/2023).
- [31] Bluetooth. URL: <https://www.bluetooth.com> (visited on 10/29/2023).
- [32] LoRa Alliance. URL: <https://loro-alliance.org> (visited on 10/29/2023).
- [33] Thread. URL: <https://www.threadgroup.org> (visited on 10/29/2023).
- [34] Connectivity Standards Alliance - Matter. URL: <https://csa-iot.org/all-solutions/matter/> (visited on 10/29/2023).
- [35] Christos Stergiou et al. "Secure integration of IoT and Cloud Computing". In: *Future Generation Computer Systems* (2018), pp. 964–975. ISSN: 0167-739X. DOI: [10.1016/j.future.2016.11.031](https://doi.org/10.1016/j.future.2016.11.031).
- [36] Moataz Soliman et al. "Smart Home: Integrating Internet of Things with Web Services and Cloud Computing". en. In: *2013 IEEE 5th International Conference on Cloud Computing Technology and Science*. 2013, pp. 317–320. DOI: [10.1109/CloudCom.2013.155](https://doi.org/10.1109/CloudCom.2013.155).

-
- [37] MQTT. The Standard for IoT Messaging. URL: <https://mqtt.org> (visited on 10/29/2023).
- [38] Zach Shelby, Klaus Hartke, and Carsten Bormann. *The Constrained Application Protocol (CoAP)*. RFC 7252. 2014. DOI: [10.17487/RFC7252](https://doi.org/10.17487/RFC7252). URL: <https://datatracker.ietf.org/doc/rfc7252>.
- [39] Jiehan Zhou et al. "CloudThings: A common architecture for integrating the Internet of Things with Cloud Computing". en. In: CSCWD, 2013, pp. 651–657. DOI: [10.1109/CSCWD.2013.6581037](https://doi.org/10.1109/CSCWD.2013.6581037).
- [40] Matthias Kovatsch, Martin Lanter, and Simon Duquennoy. "Actinium: A RESTful runtime container for scriptable Internet of Things applications". In: *2012 3rd IEEE International Conference on the Internet of Things*. 2012, pp. 135–142. DOI: [10.1109/IOT.2012.6402315](https://doi.org/10.1109/IOT.2012.6402315).
- [41] Sergio Martin et al. *Arm Mbed – AWS IoT System Integration*. en. Research Reports or Papers. 2018. URL: <https://eprints.gla.ac.uk/157277/>.
- [42] *AWS IoT Core - MQTT*. URL: <https://docs.aws.amazon.com/iot/latest/developerguide/mqtt.html> (visited on 10/29/2023).
- [43] Euripides G. M. Petrakis et al. "Internet of Things as a Service (iTaaS): Challenges and solutions for management of sensor data on the cloud and the fog". In: *Internet of Things* (2018), pp. 156–174. ISSN: 2542-6605. DOI: [10.1016/j.iot.2018.09.009](https://doi.org/10.1016/j.iot.2018.09.009).
- [44] Maria Júlia Berriel de Sousa et al. "Over-the-air firmware update for IoT devices on the wild". In: *Internet of Things* (2022). ISSN: 2542-6605. DOI: [10.1016/j.iot.2022.100578](https://doi.org/10.1016/j.iot.2022.100578).
- [45] *Amazon Elastic Load Balancer*. URL: <https://aws.amazon.com/en/elasticloadbalancing/> (visited on 10/30/2023).
- [46] Yeongho Choi and Yujin Lim. "Optimization Approach for Resource Allocation on Cloud Computing for IoT". en. In: *International Journal of Distributed Sensor Networks* (2016). ISSN: 1550-1329. DOI: [10.1155/2016/3479247](https://doi.org/10.1155/2016/3479247).
- [47] *Comment établir la coexistence entre IPv4 et IPv6*. URL: <https://community.fs.com/fr/article/how-to-achieve-ipv4-and-ipv6-coexistence-dual-stack-or-mpls-tunnel.html> (visited on 10/30/2023).
- [48] *Update the AWS IoT Greengrass Core software (OTA)*. URL: <https://docs.aws.amazon.com/greengrass/v2/developerguide/update-greengrass-core-v2.html> (visited on 10/30/2023).
- [49] *Check Greengrass core device status*. URL: <https://docs.aws.amazon.com/greengrass/v2/developerguide/device-status.html> (visited on 10/30/2023).
- [50] *What is Device Update for IoT Hub?* URL: <https://learn.microsoft.com/en-us/azure/iot-hub-device-update/understand-device-update> (visited on 10/30/2023).
- [51] *Monitor device connection status*. URL: <https://learn.microsoft.com/en-us/azure/iot-hub/monitor-device-connection-state> (visited on 10/30/2023).

Bibliography

- [52] Brendan Moran et al. *A Firmware Update Architecture for Internet of Things*. 2021. DOI: [10.17487/RFC9019](https://doi.org/10.17487/RFC9019).
- [53] Alessio Botta et al. "Integration of Cloud computing and Internet of Things: A survey". In: *Future Generation Computer Systems* (2016), pp. 684–700. ISSN: 0167-739X. DOI: [10.1016/j.future.2015.09.021](https://doi.org/10.1016/j.future.2015.09.021).
- [54] Mohammad Aazam and Eui-Nam Huh. "Fog Computing and Smart Gateway Based Communication for Cloud of Things". In: *2014 International Conference on Future Internet of Things and Cloud*. 2014, pp. 464–470. DOI: [10.1109/FiCloud.2014.83](https://doi.org/10.1109/FiCloud.2014.83).
- [55] Muhammad Ilyas, Muneeb Ahmad, and Sajid Saleem. "Internet-of-Things-Infrastructure-as-a-Service: The democratization of access to public Internet-of-Things Infrastructure". en. In: *International Journal of Communication Systems* (2020). DOI: [10.1002/dac.4562](https://doi.org/10.1002/dac.4562).
- [56] *Raspberry Pi*. URL: <https://www.raspberrypi.com> (visited on 10/20/2023).
- [57] Fung Po Tso et al. "The Glasgow Raspberry Pi Cloud: A Scale Model for Cloud Computing Infrastructures". In: *2013 IEEE 33rd International Conference on Distributed Computing Systems Workshops*. 2013, pp. 108–112. DOI: [10.1109/ICDCSW.2013.25](https://doi.org/10.1109/ICDCSW.2013.25).
- [58] Pekka Abrahamsson et al. "Affordable and Energy-Efficient Cloud Computing Clusters: The Bolzano Raspberry Pi Cloud Cluster Experiment". en. In: IEEE, 2013, pp. 170–175. ISBN: 978-0-7695-5095-4. DOI: [10.1109/CloudCom.2013.121](https://doi.org/10.1109/CloudCom.2013.121).
- [59] Ranjana Sikarwar, Pradeep Yadav, and Aditya Dubey. "A Survey on IOT enabled cloud platforms". In: *2020 IEEE 9th International Conference on Communication Systems and Network Technologies (CSNT)*. 2020, pp. 120–124. DOI: [10.1109/CSNT48778.2020.9115735](https://doi.org/10.1109/CSNT48778.2020.9115735).
- [60] *Gartner - Cloud Infrastructure and Platform Services Reviews and Ratings*. URL: <https://www.gartner.com/reviews/market/cloud-infrastructure-and-platform-services> (visited on 10/26/2023).
- [61] *Amazon Web Services*. URL: <https://aws.amazon.com/> (visited on 10/23/2023).
- [62] Yohanes Yohanie Fridelin Panduman, Sritrusta Sukaridhoto, and Anang Tjahjono. "A Survey of IoT Platform Comparison for Building Cyber-Physical System Architecture". In: *2019 International Seminar on Research of Information Technology and Intelligent Systems (ISRITI)*. 2019, pp. 238–243. DOI: [10.1109/ISRITI48646.2019.9034650](https://doi.org/10.1109/ISRITI48646.2019.9034650).
- [63] Alexandra A. Kirsanova, Gleb I. Radchenko, and Andrei N. Tchernykh. "Fog computing state of the art: concept and classification of platforms to support distributed computing systems". In: (2021). DOI: [10.48550/arXiv.2106.11726](https://doi.org/10.48550/arXiv.2106.11726).
- [64] *Microsoft Azure*. URL: <https://azure.microsoft.com/> (visited on 10/23/2023).
- [65] *Google Cloud*. URL: <https://cloud.google.com/> (visited on 10/23/2023).
- [66] *Akenza AG - Google IoT Core is Shutting Down*. URL: <https://blog.akenza.io/google-iot-core-shutdown> (visited on 10/25/2023).

- [67] Claus Pahl and Brian Lee. "Containers and Clusters for Edge Cloud Architectures – A Technology Review". In: *2015 3rd International Conference on Future Internet of Things and Cloud*. 2015, pp. 379–386. DOI: [10.1109/FiCloud.2015.35](https://doi.org/10.1109/FiCloud.2015.35).
- [68] Matti Valkeinen. "Cloud Infrastructure Tools For Cloud Applications: Infrastructure management of multiple cloud platforms". en. In: (2022). URL: <https://trepo.tuni.fi/handle/10024/137424>.
- [69] *AWS CloudFormation*. URL: <https://aws.amazon.com/en/cloudformation/> (visited on 10/26/2023).
- [70] *AWS CDK*. URL: <https://aws.amazon.com/en/cdk/> (visited on 10/26/2023).
- [71] *Azure Resource Manager*. URL: <https://azure.microsoft.com/en-gb/get-started/azure-portal/resource-manager> (visited on 10/26/2023).
- [72] *Terraform*. URL: <https://www.terraform.io> (visited on 10/26/2023).
- [73] *Pulumi*. URL: <https://www.pulumi.com> (visited on 10/26/2023).
- [74] *GitHub - hashicorp/terraform*. URL: <https://github.com/hashicorp/terraform/tags> (visited on 10/26/2023).
- [75] *GitHub - pulumi/pulumi*. URL: <https://github.com/pulumi/pulumi> (visited on 10/26/2023).
- [76] *SystemReady Certification Program*. URL: <https://www.arm.com/architecture/system-architectures/systemready-certification-program> (visited on 10/27/2023).
- [77] *Google, Oracle's Ampere VMs get Arm's SystemReady seal of approval*. URL: https://www.theregister.com/2022/11/02/google_oracles_cloud_arm/ (visited on 10/27/2023).
- [78] *SystemReady Pre-Silicon Reference Guide BSA integration and compliance*. BSA and SBSA overview. URL: <https://developer.arm.com/documentation/102858/0100/BSA-and-SBSA-overview> (visited on 10/27/2023).
- [79] *GitHub - ARM-software/bbr-ac*. URL: <https://github.com/ARM-software/bbr-ac> (visited on 10/27/2023).

Glossary

Arm Arm is a British company that designs processors with RISC-type (Reduced Instruction Set Computer) architectures. [viii](#), [3](#), [15](#), [18](#), [19](#), [24](#), [26](#), [31](#), [32](#)

AWS AWS stands for Amazon Web Services. It is a company that provides a cloud platform with over 200 services. The resources are spread over more than 32 regions around the world. [x](#), [1](#), [3](#), [20](#), [24](#), [25](#), [26](#), [28](#), [29](#), [30](#), [31](#)

cloud The term cloud refers to a remote infrastructure made up of servers that can be accessed from anywhere in the world via an internet connection. [v](#), [viii](#), [x](#), [1](#), [2](#), [3](#), [15](#), [16](#), [17](#), [18](#), [19](#), [20](#), [21](#), [22](#), [23](#), [24](#), [25](#), [26](#), [27](#), [28](#), [29](#), [30](#), [31](#), [32](#)

cloud infrastructure A cloud infrastructure is the set of computing resources in the cloud that make up an IT environment. This includes servers, storage, network and security. [x](#), [1](#), [2](#), [3](#), [16](#), [17](#), [21](#), [26](#), [27](#), [29](#), [30](#)

HES-SO The HES-SO is a Swiss university based in French-speaking Switzerland. It offers a wide range of specialisations in different fields such as technology, health, social work, economics and art. It has several buildings spread across the French-speaking cantons. [8](#)

KanBan KanBan is a project management tool designed to efficiently manage a workflow over a given period. It is widely adopted in agile methodology. [x](#), [4](#), [7](#), [10](#), [11](#)

provisioning In this project, provisioning means the automated linking of an embedded system with a cloud infrastructure. This includes the creation of a digital twin of the device, enabling the device to communicate with cloud services.. [3](#), [27](#)

Scrum Scrum is a process framework based on agile methodology for software engineering. It is suitable for small development teams. It ensures good communication between the customer and the team. It enables the first versions of a software product to be delivered quickly.. [x](#), [4](#), [7](#), [8](#), [9](#), [10](#), [11](#)

work package In project management, work packages are commonly used to define a plan. The plan provides an overview of the progress of the project. [3](#), [4](#), [7](#), [13](#)

Acronyms

API Application Programming Interface. [26](#), [29](#), [31](#)

ARM Azure Resource Manager. [29](#), [31](#)

BBR Base Boot Requirements. [31](#)

BLE Bluetooth Low Energy. [23](#)

BSA Base System Architecture. [31](#)

BSD Berkeley Software Distribution. [32](#)

CD Continuous Delivery. [3](#), [7](#), [12](#), [13](#)

CDK Cloud Development Kit. [x](#), [28](#), [31](#)

CI Continuous Intergration. [3](#), [7](#), [12](#), [13](#)

CLI Command Line interface. [30](#)

CoAP Constrained Application Protocol. [24](#)

DevOps Development and Operations. [x](#), [2](#), [12](#)

EC2 Elastic Compute Cloud. [20](#)

HTTP HyperText Transfer Protocol. [24](#)

IaaS Infrastructure as a Service. [16](#), [21](#), [29](#)

IaC Infrastructure as Code. [15](#), [18](#), [19](#), [27](#), [28](#), [29](#), [30](#), [31](#)

IoT Internet of Things. [viii](#), [x](#), [1](#), [2](#), [15](#), [18](#), [19](#), [21](#), [22](#), [23](#), [24](#), [25](#), [26](#), [27](#), [28](#), [32](#)

JSON JavaScript Object Notation. [24](#), [29](#)

MQTT Message Queuing Telemetry Transport. [24](#)

NAS Network Attached Storage. [26](#)

NAT Network Address Translation. [24](#)

Acronyms

OS Operating System. [31](#), [32](#)

OTA Over-The-Air. [25](#)

PaaS Platform as a Service. [16](#), [21](#), [27](#), [29](#)

PO Product Owner. [9](#), [10](#)

QoS Quality of Service. [24](#), [25](#)

REST REpresentational State Transfer. [24](#)

ROI Return On OInvestment. [9](#)

SaaS Software as a Service. [16](#), [29](#)

SDK Software Development Kit. [31](#)

SOA Service-Oriented Architecture. [24](#)

TCP Transmission Control Protocol. [24](#)

UDP User Datagram Protocol. [24](#)

WiFi Wireless Fidelity. [23](#)

WIP Work In Progress. [10](#), [11](#)